

Thyroid Cancer Risk Factors and Prognosis

Type Description:

This dataset contains medical and lifestyle factors associated with thyroid cancer risk, diagnosis, and prognosis. It includes patient demographics, clinical history, genetic predispositions, lifestyle habits, tumor characteristics, treatment details, and survival outcomes.

The dataset is structured to support predictive modeling, statistical analysis, and machine learning applications for risk assessment, early detection, and treatment outcome predictions.

Potential Features:

Patient Information: Age, Gender, BMI

Medical History: Family History of Thyroid Cancer, Previous Cancers, Autoimmune Disorders (e.g., Hashimoto's Thyroiditis)

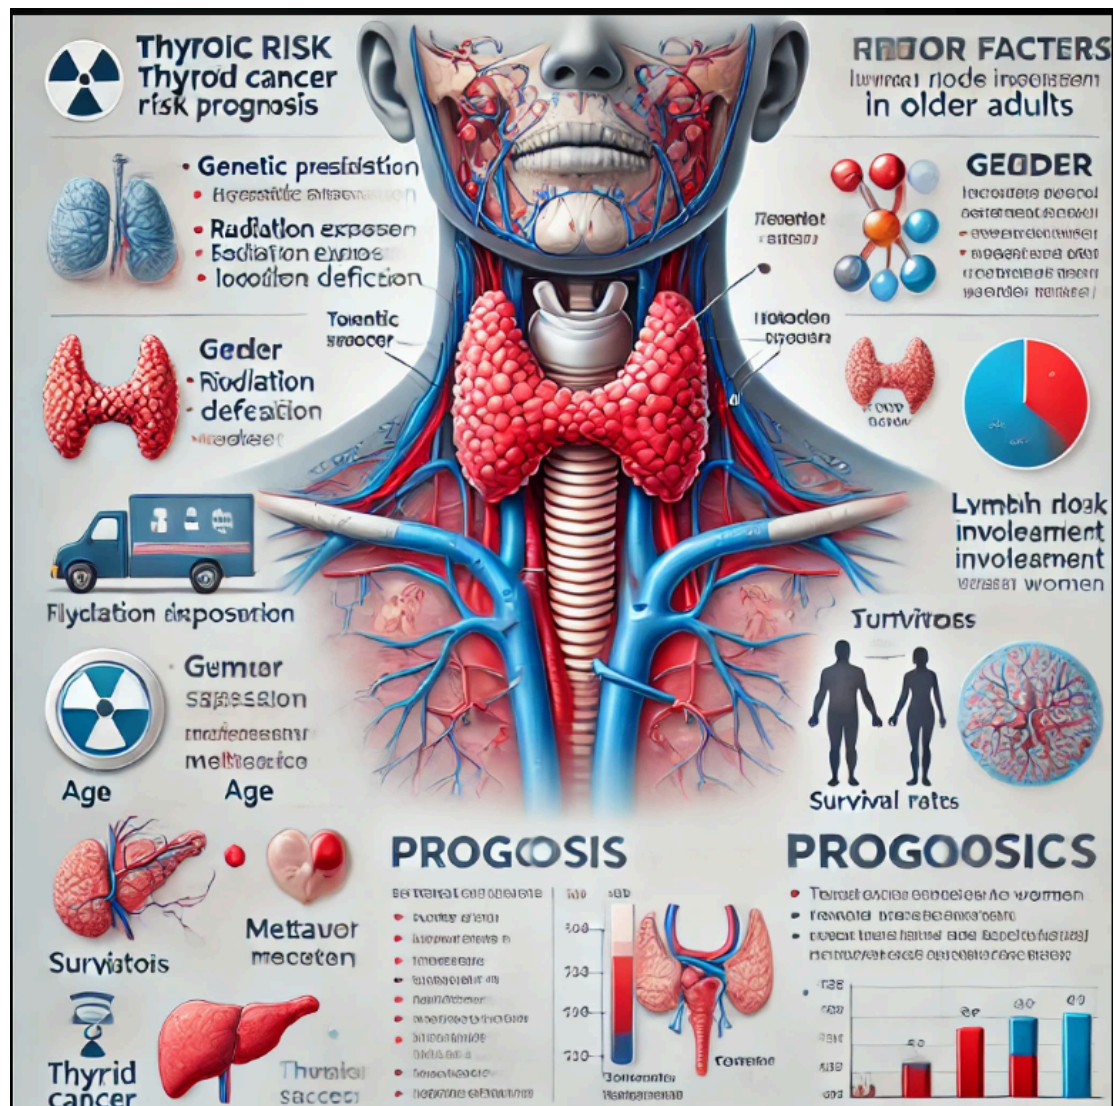
Genetic Factors: Presence of RET/PTC, BRAF, and other relevant mutations

Lifestyle Factors: Smoking Status, Alcohol Consumption, Iodine Intake, Radiation Exposure

Clinical Presentation: Tumor Size, Tumor Type (Papillary, Follicular, Medullary, Anaplastic), Nodule Characteristics, Symptom Severity

Diagnosis Details: Fine-Needle Aspiration Biopsy (FNAB) Results, Ultrasound Findings, Diagnosis Delay in Days

Treatment & Prognosis: Surgery Type, Radioactive Iodine Therapy, Chemotherapy, Recurrence Rate, Survival Years After Diagnosis



Import Libraries

```
In [5]: import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

Import Dataset

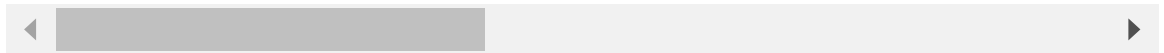
```
In [8]: df = pd.read_csv(r"C:\Users\chitt\Downloads\thyroid_cancer_risk_data.csv")
```

```
In [10]: df
```

Out[10]:

	Patient_ID	Age	Gender	Country	Ethnicity	Family_History	Radiation_Exposure
0	1	66	Male	Russia	Caucasian	No	No
1	2	29	Male	Germany	Hispanic	No	No
2	3	86	Male	Nigeria	Caucasian	No	No
3	4	75	Female	India	Asian	No	No
4	5	35	Female	Germany	African	Yes	No
...
212686	212687	58	Female	India	Asian	No	No
212687	212688	89	Male	Japan	Middle Eastern	No	No
212688	212689	72	Female	Nigeria	Hispanic	No	No
212689	212690	85	Female	Brazil	Middle Eastern	No	No
212690	212691	46	Female	Japan	Middle Eastern	No	No

212691 rows × 17 columns

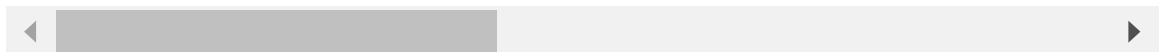


In [12]: df.shape

Out[12]: (212691, 17)

In [14]: df.head()

	Patient_ID	Age	Gender	Country	Ethnicity	Family_History	Radiation_Exposure	Is_Thyroid_Cancer
0	1	66	Male	Russia	Caucasian	No	No	Yes
1	2	29	Male	Germany	Hispanic	No	No	Yes
2	3	86	Male	Nigeria	Caucasian	No	No	No
3	4	75	Female	India	Asian	No	No	No
4	5	35	Female	Germany	African	Yes	Yes	Yes



In [16]: df.tail()

Out[16]:

	Patient_ID	Age	Gender	Country	Ethnicity	Family_History	Radiation_Exposure
212686	212687	58	Female	India	Asian	No	No
212687	212688	89	Male	Japan	Middle Eastern	No	No
212688	212689	72	Female	Nigeria	Hispanic	No	No
212689	212690	85	Female	Brazil	Middle Eastern	No	No
212690	212691	46	Female	Japan	Middle Eastern	No	No

In [20]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 212691 entries, 0 to 212690
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Patient_ID            212691 non-null int64
1   Age                   212691 non-null int64
2   Gender                212691 non-null object
3   Country               212691 non-null object
4   Ethnicity             212691 non-null object
5   Family_History        212691 non-null object
6   Radiation_Exposure    212691 non-null object
7   Iodine_Deficiency     212691 non-null object
8   Smoking               212691 non-null object
9   Obesity               212691 non-null object
10  Diabetes              212691 non-null object
11  TSH_Level             212691 non-null float64
12  T3_Level              212691 non-null float64
13  T4_Level              212691 non-null float64
14  Nodule_Size           212691 non-null float64
15  Thyroid_Cancer_Risk   212691 non-null object
16  Diagnosis             212691 non-null object
dtypes: float64(4), int64(2), object(11)
memory usage: 27.6+ MB
```

In [22]: `df.describe()`

Out[22]:

	Patient_ID	Age	TSH_Level	T3_Level	T4_Level	No
count	212691.00000	212691.000000	212691.000000	212691.000000	212691.000000	212691.000000
mean	106346.00000	51.918497	5.045102	2.001727	8.246204	
std	61398.74739	21.632815	2.860264	0.866248	2.164188	
min	1.00000	15.000000	0.100000	0.500000	4.500000	
25%	53173.50000	33.000000	2.570000	1.250000	6.370000	
50%	106346.00000	52.000000	5.040000	2.000000	8.240000	
75%	159518.50000	71.000000	7.520000	2.750000	10.120000	
max	212691.00000	89.000000	10.000000	3.500000	12.000000	

In [24]: df.corr

```

Out[24]: <bound method DataFrame.corr of
Ethnicity Family_History \
0          1    66    Male    Russia    Caucasian    No
1          2    29    Male    Germany    Hispanic    No
2          3    86    Male    Nigeria    Caucasian    No
3          4    75    Female    India    Asian    No
4          5    35    Female    Germany    African    Yes
...      ...    ...    ...    ...    ...    ...
212686    212687    58    Female    India    Asian    No
212687    212688    89    Male    Japan    Middle Eastern    No
212688    212689    72    Female    Nigeria    Hispanic    No
212689    212690    85    Female    Brazil    Middle Eastern    No
212690    212691    46    Female    Japan    Middle Eastern    No

Radiation_Exposure Iodine_Deficiency Smoking Obesity Diabetes \
0          Yes    No    No    No    No
1          Yes    No    No    No    No
2          No    No    No    No    No
3          No    No    No    No    No
4          Yes    No    No    No    No
...      ...    ...    ...    ...    ...
212686    No    No    No    Yes    No
212687    No    No    No    Yes    No
212688    No    No    No    No    Yes
212689    No    No    No    No    Yes
212690    No    No    Yes    No    No

TSH_Level T3_Level T4_Level Nodule_Size Thyroid_Cancer_Risk \
0      9.37    1.67    6.16    1.08    Low
1      1.83    1.73    10.54    4.05    Low
2      6.26    2.59    10.57    4.61    Low
3      4.10    2.62    11.04    2.46    Medium
4      9.10    2.11    10.71    2.11    High
...      ...    ...    ...    ...    ...
212686    2.00    0.64    11.92    1.48    Low
212687    9.77    3.25    7.30    4.46    Medium
212688    7.72    2.44    8.71    2.36    Medium
212689    5.62    2.53    9.62    1.54    Medium
212690    5.60    2.73    10.59    2.53    Low

Diagnosis
0      Benign
1      Benign
2      Benign
3      Benign
4      Benign
...      ...
212686    Benign
212687    Benign
212688    Benign
212689    Benign
212690    Malignant

[212691 rows x 17 columns]>

```

```
In [26]: df.isnull().sum()
```

```
Out[26]: Patient_ID      0
        Age             0
        Gender          0
        Country         0
        Ethnicity       0
        Family_History  0
        Radiation_Exposure 0
        Iodine_Deficiency 0
        Smoking         0
        Obesity         0
        Diabetes        0
        TSH_Level       0
        T3_Level        0
        T4_Level        0
        Nodule_Size     0
        Thyroid_Cancer_Risk 0
        Diagnosis       0
        dtype: int64
```

```
In [28]: df.duplicated()
```

```
Out[28]: 0      False
        1      False
        2      False
        3      False
        4      False
        ...
        212686 False
        212687 False
        212688 False
        212689 False
        212690 False
        Length: 212691, dtype: bool
```

```
In [30]: df.columns
```

```
Out[30]: Index(['Patient_ID', 'Age', 'Gender', 'Country', 'Ethnicity', 'Family_History',
               'Radiation_Exposure', 'Iodine_Deficiency', 'Smoking', 'Obesity',
               'Diabetes', 'TSH_Level', 'T3_Level', 'T4_Level', 'Nodule_Size',
               'Thyroid_Cancer_Risk', 'Diagnosis'],
              dtype='object')
```

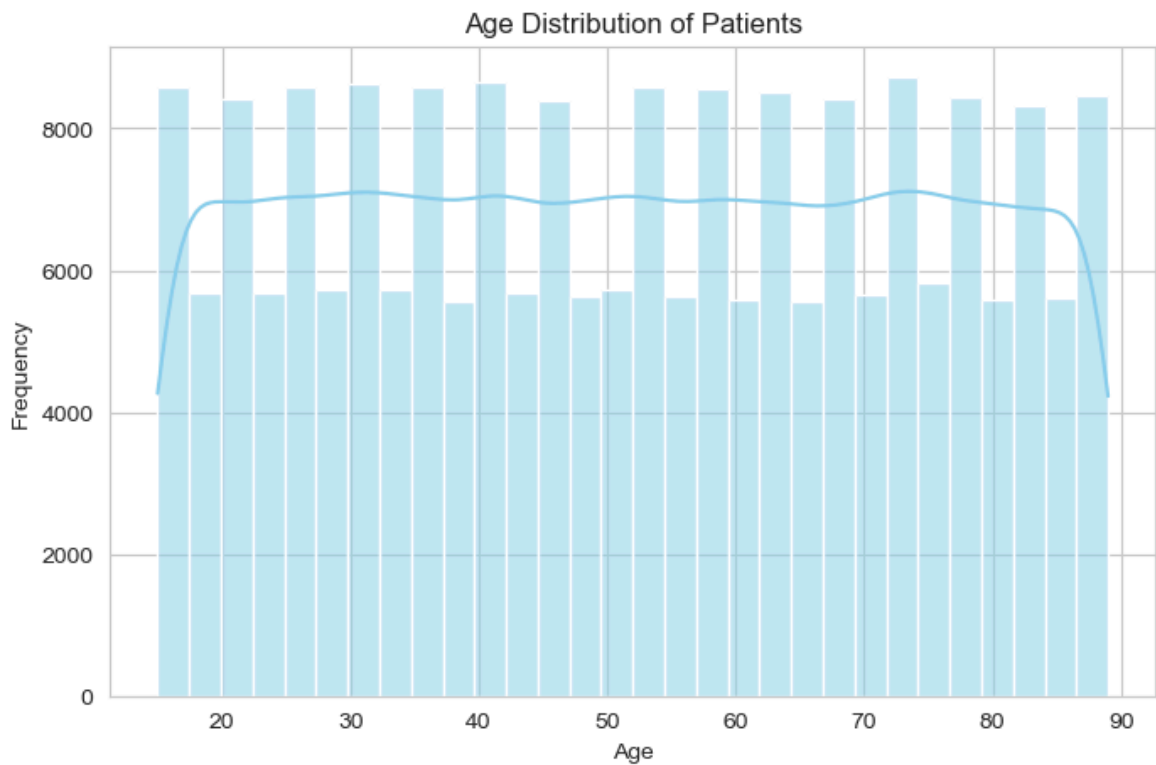
```
In [32]: df["Family_History"] = df["Family_History"].map({"Yes": 1, "No": 0}) # Adjust b
df["Thyroid_Cancer_Risk"] = pd.to_numeric(df["Thyroid_Cancer_Risk"], errors="coe
```

Data Visualizations

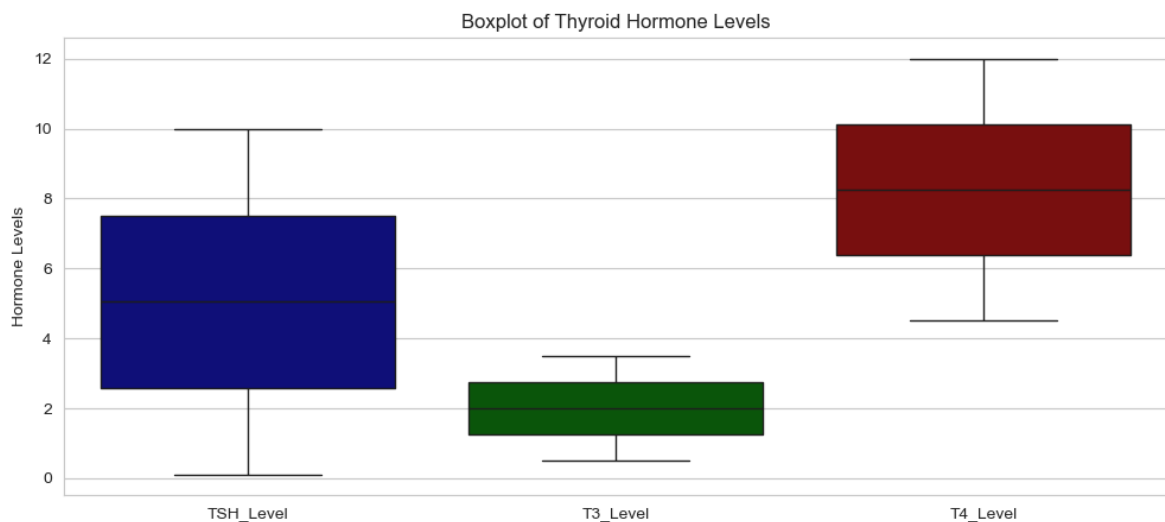
```
In [35]: # Set style
sns.set_style("whitegrid")
plt.rcParams["figure.figsize"] = (10, 6)
```

```
In [37]: # 1. Age Distribution
plt.figure(figsize=(8, 5))
sns.histplot(df["Age"], kde=True, bins=30, color="skyblue")
plt.title("Age Distribution of Patients")
plt.xlabel("Age")
```

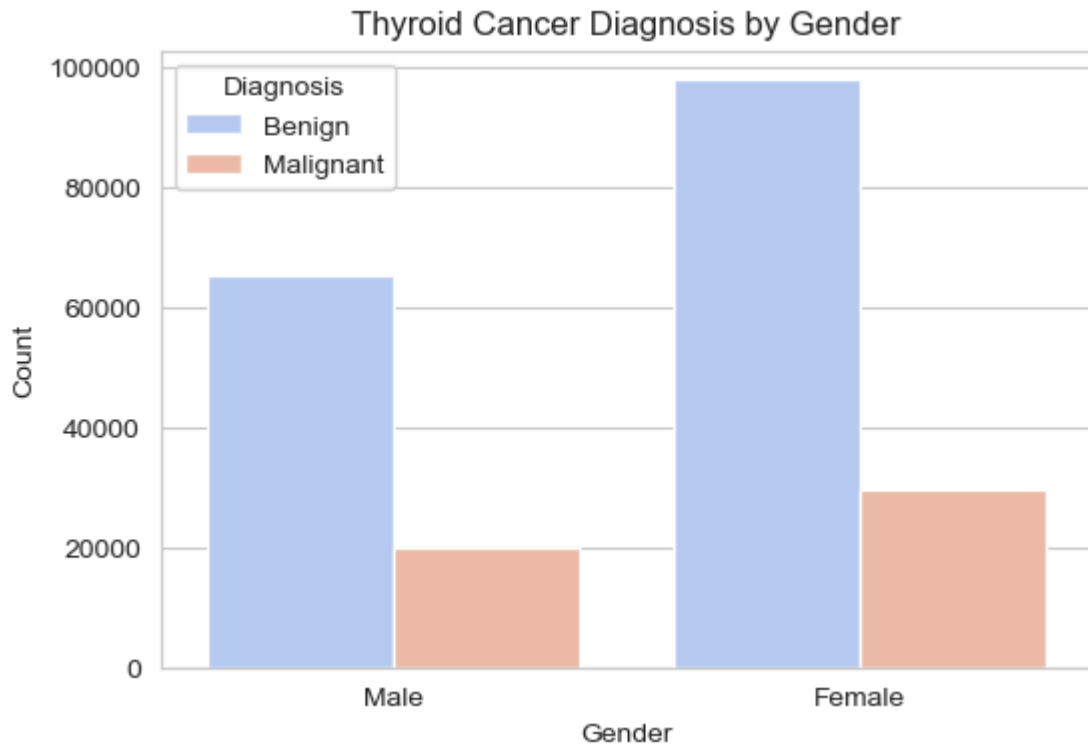
```
plt.ylabel("Frequency")
plt.show()
```



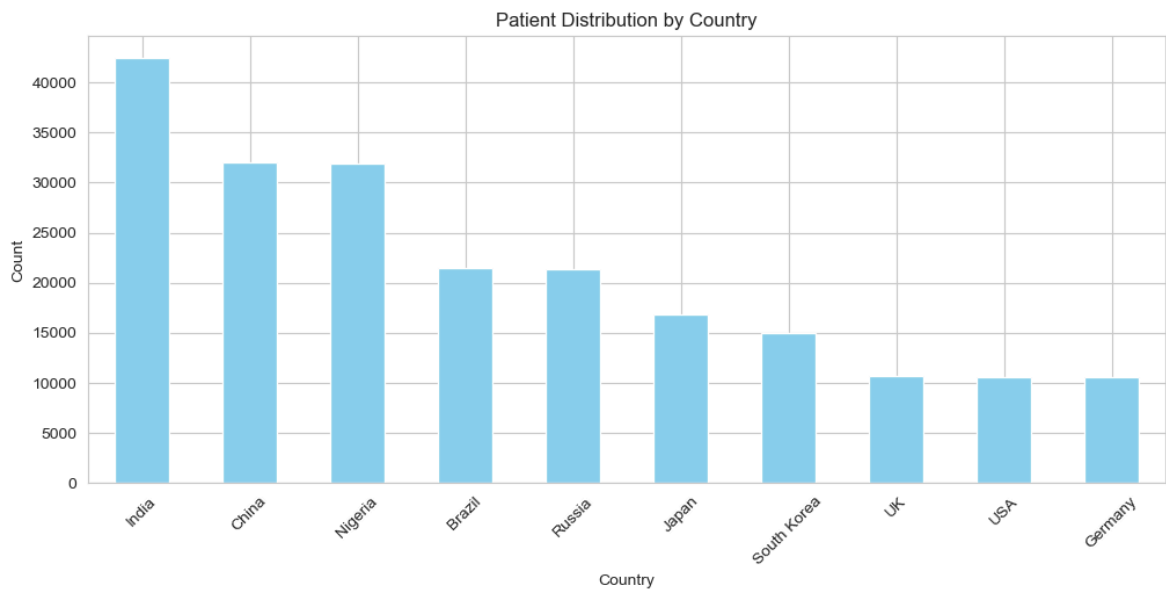
```
In [39]: # 2. Boxplot of TSH, T3, and T4 Levels
plt.figure(figsize=(12, 5))
sns.boxplot(data=df[["TSH_Level", "T3_Level", "T4_Level"]], palette=["darkblue",
plt.title("Boxplot of Thyroid Hormone Levels")
plt.ylabel("Hormone Levels")
plt.show()
```



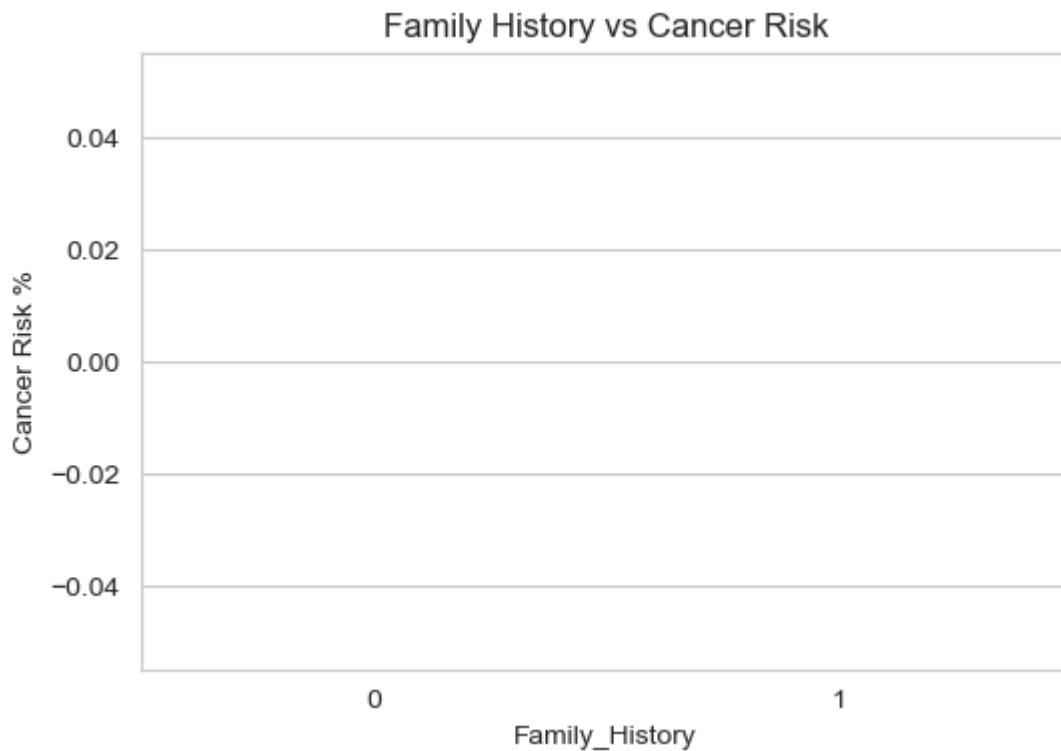
```
In [41]: # 3. Gender vs Diagnosis
plt.figure(figsize=(6, 4))
sns.countplot(x="Gender", hue="Diagnosis", data=df, palette="coolwarm")
plt.title("Thyroid Cancer Diagnosis by Gender")
plt.xlabel("Gender")
plt.ylabel("Count")
plt.show()
```

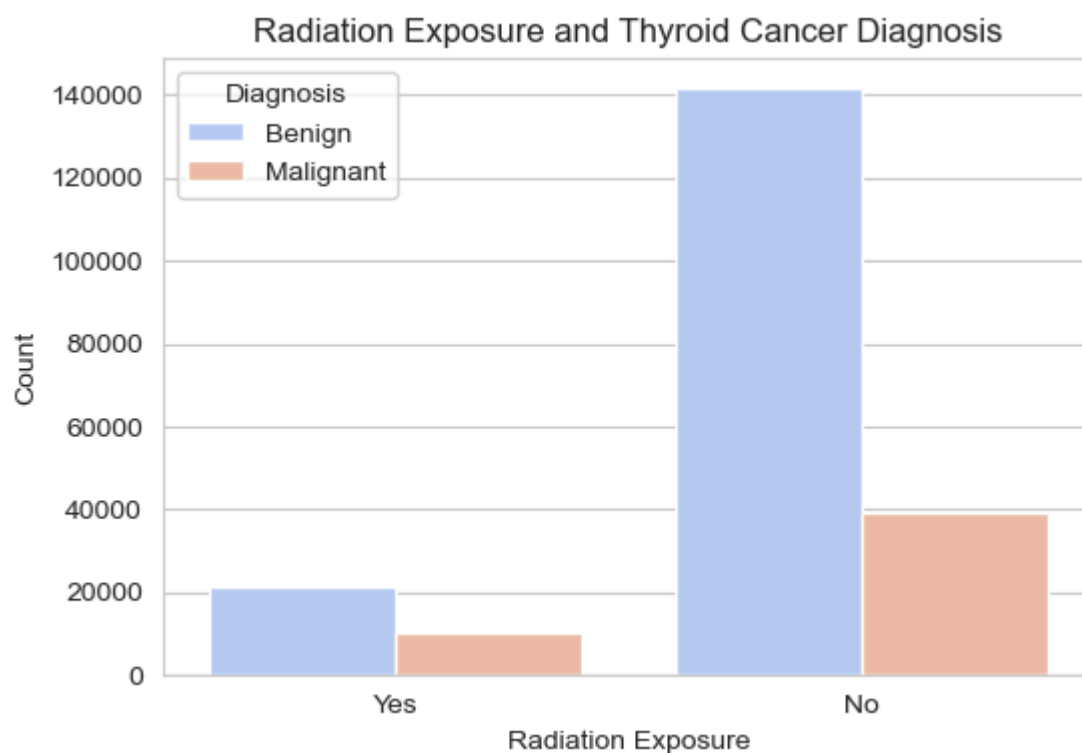
```
In [43]: # 4. Country-wise Distribution
plt.figure(figsize=(12, 5))
df["Country"].value_counts().plot(kind="bar", color="skyblue")
plt.title("Patient Distribution by Country")
plt.xlabel("Country")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()
```



```
In [45]: # 5. Family History vs Thyroid Cancer Risk
plt.figure(figsize=(6, 4))
sns.barplot(x="Family_History", y="Thyroid_Cancer_Risk", data=df, palette="Blues")
plt.title("Family History vs Cancer Risk")
plt.ylabel("Cancer Risk %")
plt.show()
```

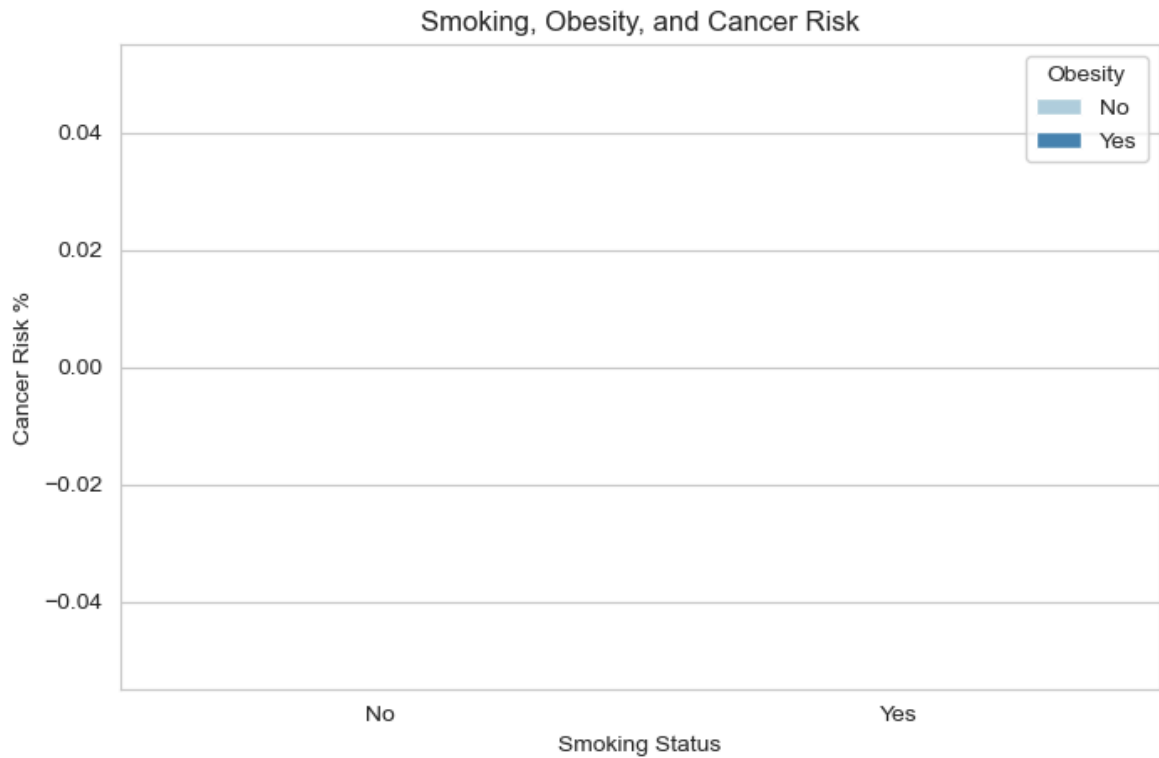


```
In [47]: # 6. Radiation Exposure & Cancer Diagnosis
plt.figure(figsize=(6, 4))
sns.countplot(x="Radiation_Exposure", hue="Diagnosis", data=df, palette="coolwarm")
plt.title("Radiation Exposure and Thyroid Cancer Diagnosis")
plt.xlabel("Radiation Exposure")
plt.ylabel("Count")
plt.show()
```

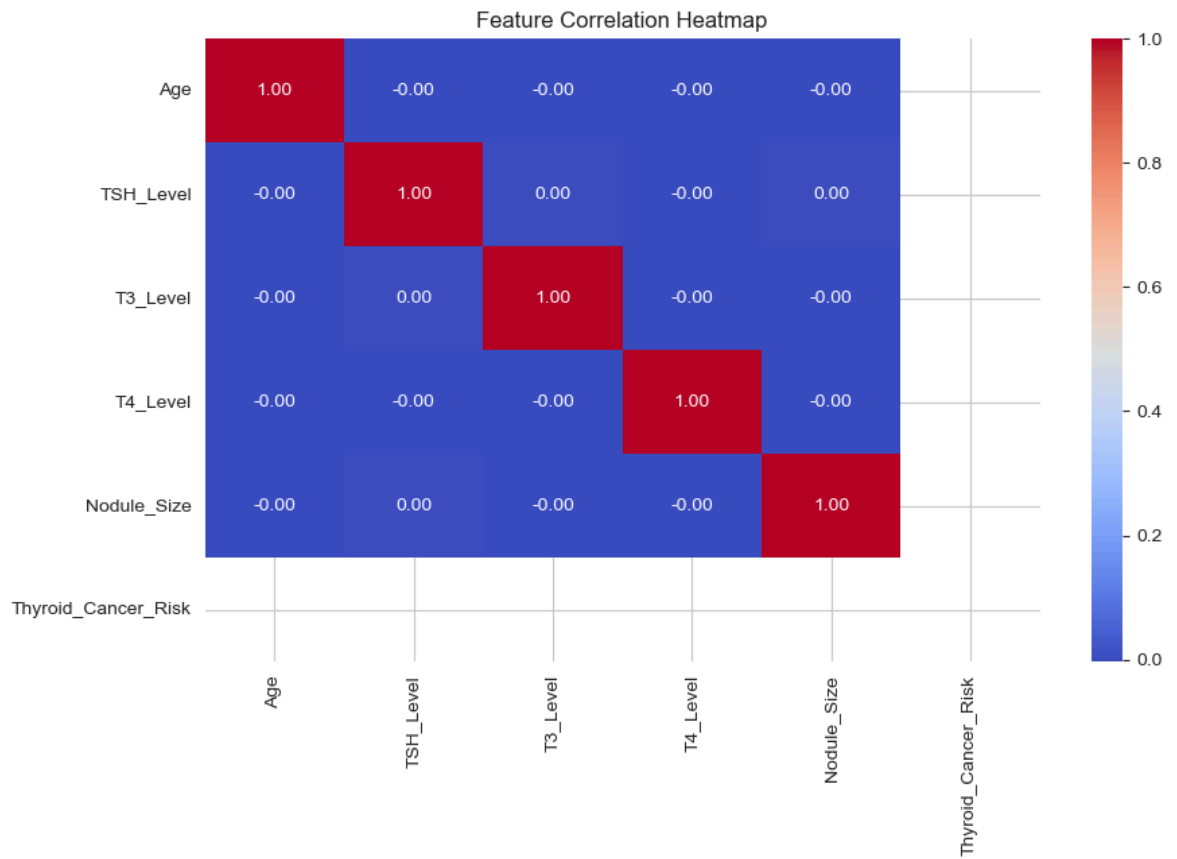


```
In [49]: # 7. Smoking & Obesity vs Cancer Risk
plt.figure(figsize=(8, 5))
sns.barplot(x="Smoking", y="Thyroid_Cancer_Risk", hue="Obesity", data=df, palette="coolwarm")
plt.title("Smoking, Obesity, and Cancer Risk")
```

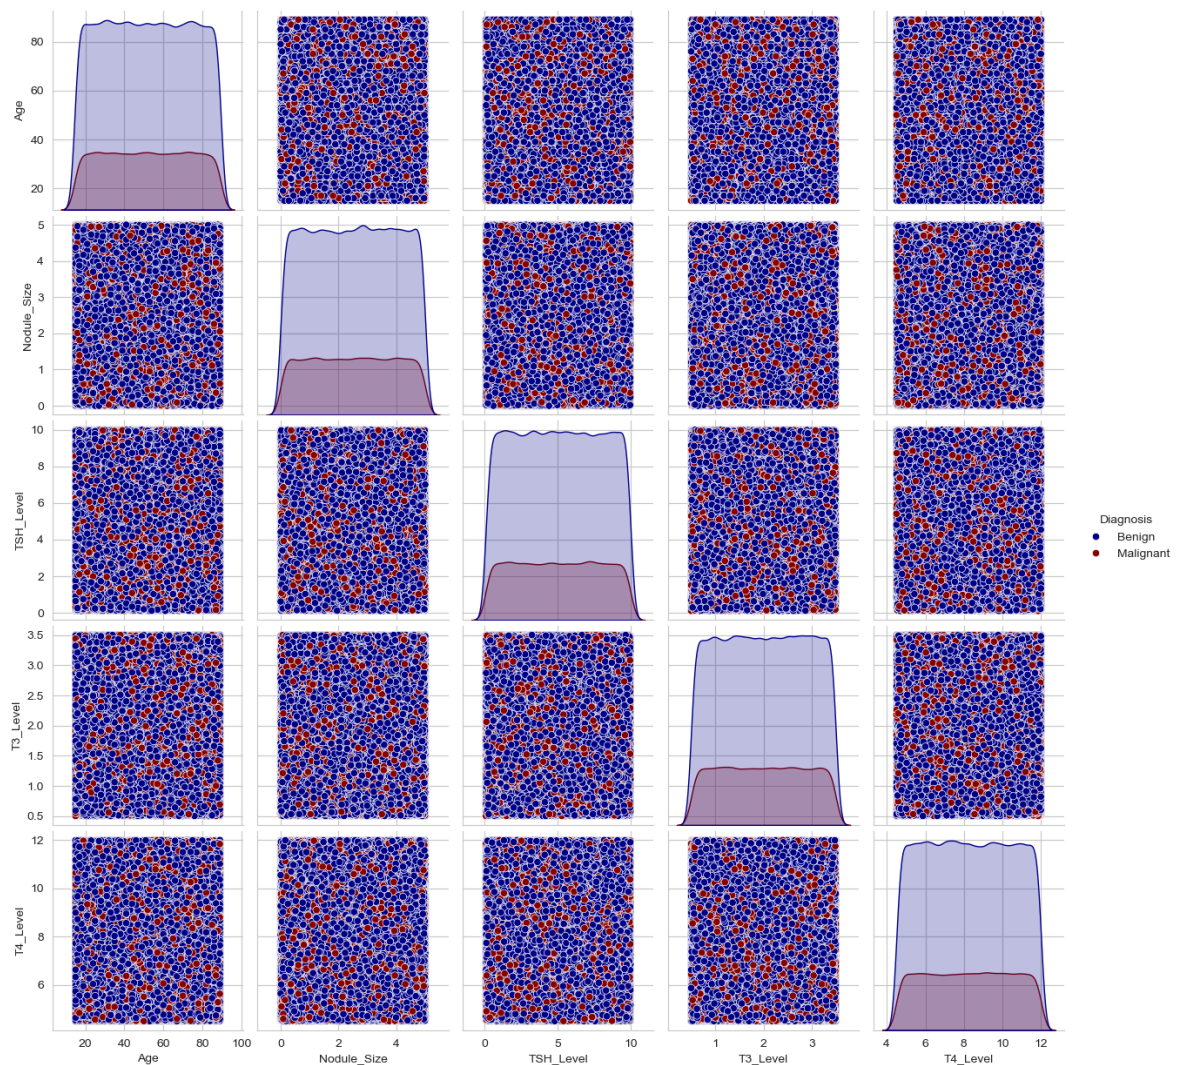
```
plt.xlabel("Smoking Status")  
plt.ylabel("Cancer Risk %")  
plt.show()
```



```
In [51]: # 8. Correlation Heatmap  
plt.figure(figsize=(10, 6))  
corr = df[["Age", "TSH_Level", "T3_Level", "T4_Level", "Nodule_Size", "Thyroid_C  
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f")  
plt.title("Feature Correlation Heatmap")  
plt.show()
```



```
In [53]: # 9. Pairplot for Key Features
sns.pairplot(df, vars=["Age", "Nodule_Size", "TSH_Level", "T3_Level", "T4_Level"])
plt.show()
```



ML Algorithms

```
In [56]: from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score, classification_report
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
```

```
In [60]: from sklearn.preprocessing import LabelEncoder
df_categorical = df.select_dtypes(include='object')
label_encoders = {}
for column in df_categorical.columns:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le
print(df.head())
```

	Patient_ID	Age	Gender	Country	Ethnicity	Family_History	\
0	1	66	1	6	2		0
1	2	29	1	2	3		0
2	3	86	1	5	2		0
3	4	75	0	3	1		0
4	5	35	0	2	0		1

	Radiation_Exposure	Iodine_Deficiency	Smoking	Obesity	Diabetes	\
0	1		0	0	0	0
1	1		0	0	0	0
2	0		0	0	0	0
3	0		0	0	0	0
4	1		0	0	0	0

	TSH_Level	T3_Level	T4_Level	Nodule_Size	Thyroid_Cancer_Risk	Diagnosis
0	9.37	1.67	6.16	1.08	NaN	0
1	1.83	1.73	10.54	4.05	NaN	0
2	6.26	2.59	10.57	4.61	NaN	0
3	4.10	2.62	11.04	2.46	NaN	0
4	9.10	2.11	10.71	2.11	NaN	0

```
In [62]: from sklearn.preprocessing import StandardScaler
```

```
x = df.drop(columns=["Patient_ID", "Diagnosis"])
y = df["Diagnosis"]
x_scaled = StandardScaler().fit_transform(x)
```

```
In [64]: from sklearn.model_selection import train_test_split
```

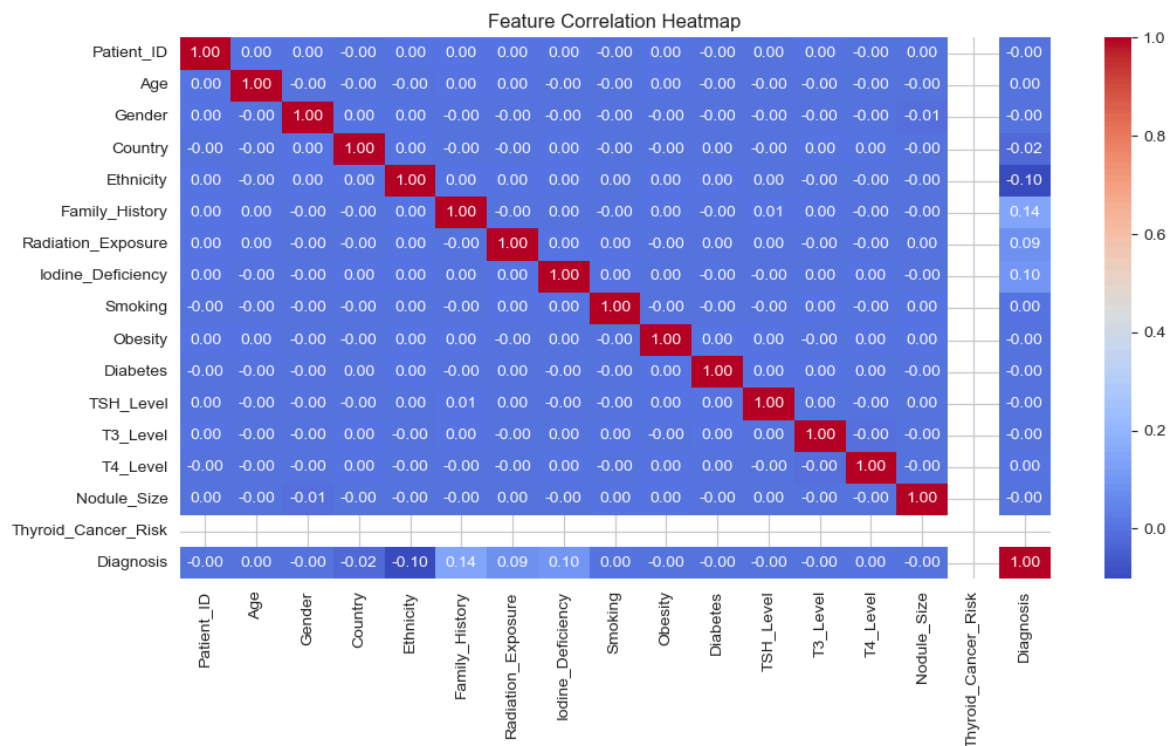
```
In [66]: x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2,
```

```
In [68]: for col in df.columns:
            if df[col].dtype == "object":
                df[col].fillna(df[col].mode()[0], inplace=True)
            else:
                df[col].fillna(df[col].median(), inplace=True)
```

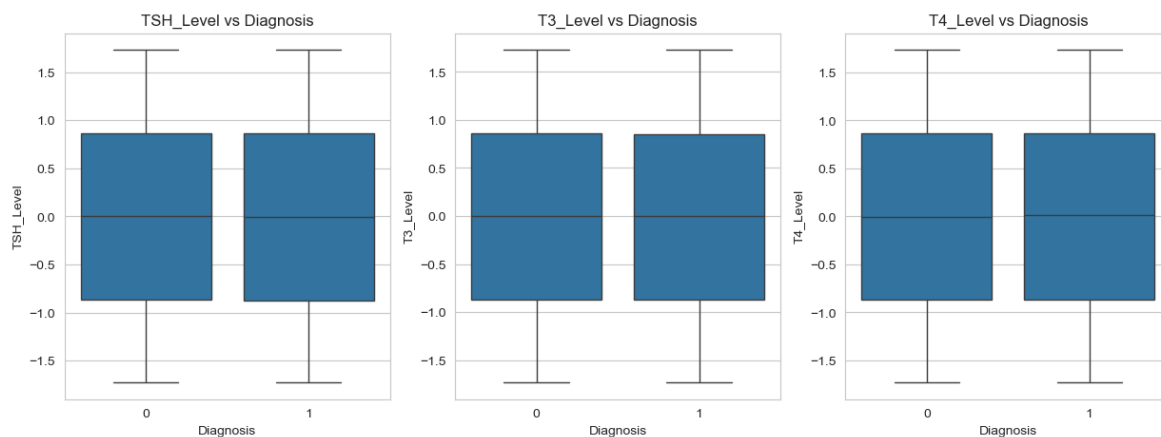
```
In [70]: categorical_cols = ['Gender', 'Country', 'Ethnicity', 'Family_History',
                             'Radiation_Exposure', 'Iodine_Deficiency', 'Smoking',
                             'Obesity', 'Diabetes']
encoder = LabelEncoder()
for col in categorical_cols:
    df[col] = encoder.fit_transform(df[col])
```

```
In [72]: scaler = StandardScaler()
num_cols = ['Age', 'TSH_Level', 'T3_Level', 'T4_Level', 'Nodule_Size']
df[num_cols] = scaler.fit_transform(df[num_cols])
```

```
In [74]: plt.figure(figsize=(12, 6))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
In [76]: plt.figure(figsize=(15, 5))
for i, col in enumerate(['TSH_Level', 'T3_Level', 'T4_Level'], 1):
    plt.subplot(1, 3, i)
    sns.boxplot(x=y, y=df[col])
    plt.title(f"{col} vs Diagnosis")
plt.show()
```



```
In [78]: print(df.isnull().sum())
```

```

Patient_ID      0
Age             0
Gender          0
Country         0
Ethnicity       0
Family_History  0
Radiation_Exposure  0
Iodine_Deficiency  0
Smoking         0
Obesity         0
Diabetes        0
TSH_Level       0
T3_Level        0
T4_Level        0
Nodule_Size     0
Thyroid_Cancer_Risk  212691
Diagnosis       0
dtype: int64

```

```
In [80]: df['Thyroid_Cancer_Risk'].fillna(df['Thyroid_Cancer_Risk'].median(), inplace=True)
```

```
In [82]: print(df.isnull().sum())
```

```

Patient_ID      0
Age             0
Gender          0
Country         0
Ethnicity       0
Family_History  0
Radiation_Exposure  0
Iodine_Deficiency  0
Smoking         0
Obesity         0
Diabetes        0
TSH_Level       0
T3_Level        0
T4_Level        0
Nodule_Size     0
Thyroid_Cancer_Risk  212691
Diagnosis       0
dtype: int64

```

```
In [84]: from sklearn.impute import SimpleImputer
imputer = SimpleImputer(strategy='mean') # or 'median', 'most_frequent'
x_train = imputer.fit_transform(x_train)
x_test = imputer.transform(x_test)
```

```
In [86]: print("Original X columns:", df.drop(columns=["Patient_ID", "Diagnosis"]).shape[1])
print("Scaled X columns:", x_scaled.shape[1])
```

```

Original X columns: 15
Scaled X columns: 15

```

```
In [88]: from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy="mean") # or "median", "most_frequent"
X_imputed = imputer.fit_transform(df.drop(columns=["Patient_ID", "Diagnosis"]))

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)
```



```
In [90]: feature_names = df.drop(columns=["Patient_ID", "Diagnosis"]).columns[:x_train.sh
x_train = pd.DataFrame(x_train, columns=feature_names)
x_test = pd.DataFrame(x_test, columns=feature_names)
```

```
In [92]: from sklearn.model_selection import RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

models_params = {
    'RandomForest': (RandomForestClassifier(random_state=42), {
        'n_estimators': [50, 100],
        'max_depth': [None, 10],
    }),
    'LogisticRegression': (LogisticRegression(random_state=42, max_iter=500), {
        'C': [0.01, 1, 100],
    }),
    'XGBoost': (XGBClassifier(random_state=42, use_label_encoder=False, eval_met
        'n_estimators': [50, 100],
        'learning_rate': [0.01, 0.1],
    ))
}
results = {}
for name, (model, param_dist) in models_params.items():
    print(f"Training {name}...")

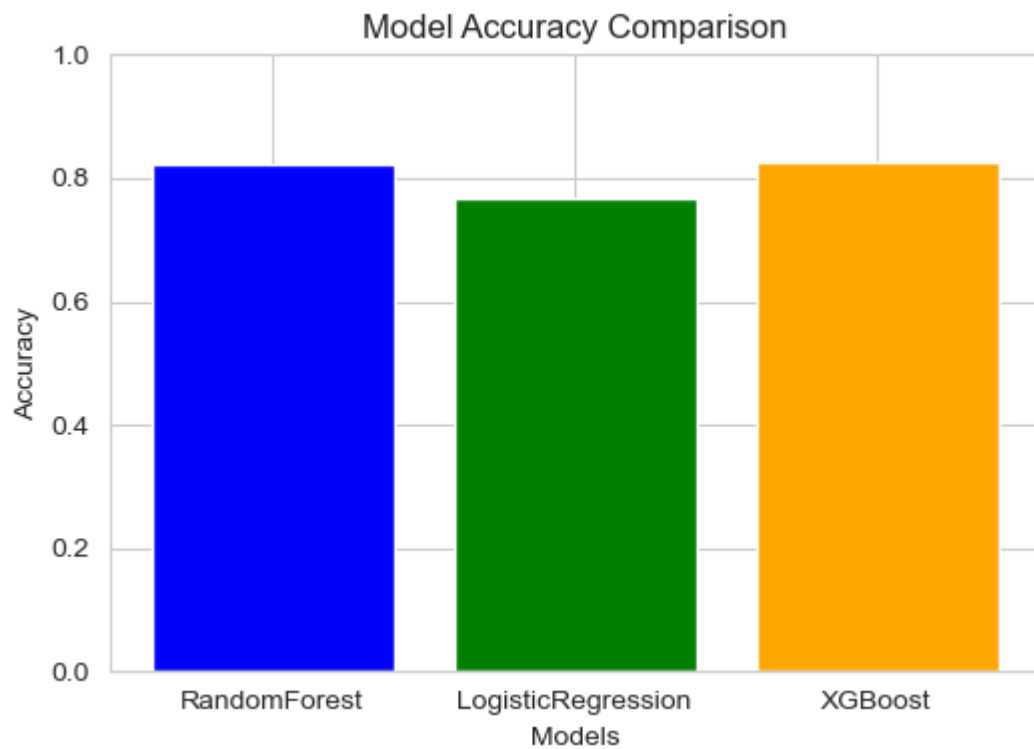
    search = RandomizedSearchCV(
        estimator=model,
        param_distributions=param_dist,
        n_iter=3,
        scoring='accuracy',
        cv=3,
        n_jobs=-1,
        random_state=42,
        verbose=0
    )
    search.fit(x_train, y_train)

    best_model = search.best_estimator_
    y_pred = best_model.predict(x_test)

    results[name] = accuracy_score(y_test, y_pred)
    print(f"{name} Accuracy: {results[name]:.4f}")

# Visualization
plt.figure(figsize=(6, 4))
plt.bar(results.keys(), results.values(), color=['blue', 'green', 'orange'])
plt.xlabel("Models")
plt.ylabel("Accuracy")
plt.title("Model Accuracy Comparison")
plt.ylim(0, 1)
plt.show()
```

Training RandomForest...
RandomForest Accuracy: 0.8228
Training LogisticRegression...
LogisticRegression Accuracy: 0.7680
Training XGBoost...
XGBoost Accuracy: 0.8249



Completed